



图形绘制技术 (Rendering)

Chapter 4: Ray Tracing Advanced

过 洁

南京大学计算机科学与技术系

guojie@nju.edu.cn




Shading the Intersection Point



- The shading of each intersection point is the sum of contributions from all light sources
 - Cast rays from the intersection point to all light sources
- Light types:
 - Ambient light, point light, directional light, spot light, area light, volume light, etc.
- Material properties:
 - Diffusion, specular, shininess, emission, etc.



A Ray Tracing Example





Ray Tracing Fails





Distributed Ray Tracing



What is distributed ray tracing?

- Distributed ray tracing is **not** ray tracing on a distributed system.
- Distributed ray tracing **is** a ray tracing method based on randomly distributed oversampling to reduce aliasing artifacts in rendered images.



Distributed Ray Tracing



■ Motivation

- The classical ray tracing produces very clean images (look fake)
 - Perfect focus
 - Perfect reflections
 - Sharp shadows

■ Main idea

- Replace the single ray with a distribution of rays

■ Add randomness to rendering

- Antialiasing
- Soft shadows
- Depth-of-field
- Motion blur
- Glossy reflections







Aliasing in Rendering



- One of the most common rendering artifacts is the “jaggies”. Consider rendering a white polygon against a black background:



- We would instead like to get a smoother transition:






Antialiasing in a ray tracer



- We would like to compute the average intensity in the neighborhood of each pixel.




- When casting one ray per pixel, we are likely to have aliasing artifacts.
- To improve matters, we can cast more than one ray per pixel and average the result.
- A.k.a., **super-sampling and averaging down**.



Antialiasing



- e.g. one sample / pixel
- for each pixel (x, y) do
 - $c(x, y) = \text{trace}(x + 0.5, y + 0.5)$







Antialiasing



- e.g. four samples / pixel
- Regular sampling: divide each pixel area to $n \times n$ grids and generate a ray within each grid
- for each pixel (x, y) do
 - $c(x, y) = 0$
 - for $p = 0$ to $n - 1$ do
 - for $q = 0$ to $n - 1$ do
 - $c(x, y) = c(x, y) + \text{trace}\left(x + \frac{p+0.5}{n}, y + \frac{q+0.5}{n}\right)$
 - $c(x, y) = c(x, y)/n^2$
- Equal to the traditional rendering
 - Uses a larger image resolution
 - Down-sample the image to make a target resolution
 - Still make the regular pattern!







Antialiasing



- e.g. four samples / pixel
- Random sampling: randomly generate n^2 rays
- for each pixel (x, y) do
 - $c(x, y) = 0$
 - for $p = 0$ to $n^2 - 1$ do
 - $c(x, y) = c(x, y) + \text{trace}(x + \epsilon, y + \epsilon)$
 - $c(x, y) = c(x, y) / n^2$
- $\epsilon \in [0, 1)$ is a random number
- The regular pattern is converted into image noise







Antialiasing





- e.g. four samples / pixel
- Jittering (stratified sampling): randomly generate a ray within each grid
- for each pixel (x, y) do
 - $c(x, y) = 0$
 - for $p = 0$ to $n - 1$ do
 - for $q = 0$ to $n - 1$ do
 - $c(x, y) = c(x, y) + \text{trace}\left(x + \frac{p+\epsilon}{n}, y + \frac{q+\epsilon}{n}\right)$
 - $c(x, y) = c(x, y)/n^2$
- This is a hybrid approach between the regular and random sampling





Soft Shadows






Soft Shadows



- How can we implement the soft shadows?
 - Approximate the area light with multiple point lights (regularly chosen)
 - Produce regular patterns
 - Distributed ray tracing: randomly choose one within the light shape






Soft Shadows





- e.g. area light defined as a parallelogram
 - Select a random point
 - $l_{pos} = c + \varepsilon_1 a + \varepsilon_2 b$
 - $\varepsilon_1, \varepsilon_2 \in [0,1)$
 - Generate a shadow ray from this point







Soft Shadows






Soft Shadows






Soft Shadows





Soft Shadows






Depth of Field



- This effect occurs when using a nonzero size “lens”
 - Approximate real cameras that have lenses with focal lengths




[Moon et al. 15]



Depth of Field




- This effect occurs when using a nonzero size “lens”
 - Approximate real cameras that have lenses with focal lengths
 - Can model the lens with a disk





DoF and Motion Blur






Gloss



- Standard ray tracing produces sharp reflections (mirror like).
- Distributed ray tracing uses multiple rays around the reflected ray to produce a blur image of the reflected object.



A Note on Material







Gloss



Sharp VS Fuzzy Gloss




Standard Ray Tracer



Distributed Ray Tracer
(50 Rays)



Distributed Ray Tracing: All in One






More General? Path Tracing



- Path tracing extends distributed ray tracing by stochastically sampling all possible light paths





Parallelization



- Ray tracing is inherently parallel, since the rays for one pixel are independent of the rays for other pixels
- Can take advantage of modern parallel CPU/GPU/Clusters to significantly accelerate a ray tracer
 - Threading (e.g., Pthread, OpenMP) distributes rays across cores
 - Message Passing Interface (MPI) distributes rays across processors on different machines
 - OptiX/CUDA distributes rays on the GPU



Parallelization



- Memory coherency helps when distributing rays to various threads/processors
 - Assign spatially neighboring rays (on the image plane) to the same core/processor
 - These rays tend to intersect with the same objects in the scene, and thus access the same memory



Real Time Ray-Tracing



- NVIDIA Optix...

